




Optimal CXL device count on H100 — throughput saturates at the HBM ceiling; optimum (orange) is N_CXL=16 for MHA long-context, N_CXL=4 for GQA



Parameters: GPU = NVIDIA H100 SXM (989 TFLOPS FP16, 3350 GB/s HBM, 80 GB) | batch = 32 | 1000 LoRA adapters (ranks {8,16,32,64,128}) | per device: 128 GB/s CXL link, 1.1 TB/s DRAM, 2 TFLOPS NDP | KV short = [100,1024] (GQA /4), long = [2048,4096] (GQA /4) tokens | seed 42, median of 3